

# Deprecations

[Copy page](#)

Find deprecated features and recommended replacements.

## Overview

As we launch safer and more capable models, we regularly retire older models. Software relying on OpenAI models may need occasional updates to keep working. Impacted customers will always be notified by email and in our documentation along with [blog posts](#) for larger changes.

This page lists all API deprecations, along with recommended replacements.

## Deprecation vs Legacy

We use the term "deprecation" to refer to the process of retiring a model or endpoint. When we announce that a model or endpoint is being deprecated, it immediately becomes deprecated. All deprecated models and endpoints will also have a shut down date. At the time of the shut down, the model or endpoint will no longer be accessible.

We use the term "legacy" to refer to models and endpoints that will no longer receive updates. We tag endpoints and models as legacy to signal to developers where we are moving as a platform and that they should likely migrate to newer models or endpoints. You can expect that a legacy model or endpoint will be deprecated at some point in the future.

## Incremental model updates

As [announced](#) in March 2023, we regularly release new versions of `gpt-4` and `gpt-3.5-turbo`.

Each model version is dated with either `-MMDD` or a `YYYY-MM-DD` suffix; e.g., `gpt-4-0613` or `gpt-4o-2024-05-13`. The undated model name, e.g., `gpt-4o`, will typically point to the latest version dated version. Users of undated model names will be notified by email typically 2 weeks before any change takes place.

## Migrating to replacements

Once a model is deprecated, be sure to migrate all usage to a suitable replacement before the shutdown date. Requests to models past the shutdown date will fail.

To help measure the performance of replacement models on your tasks, we have open-sourced [Evals](#), a Python framework for evaluating language models.

If new models perform worse on your tasks, let us know by submitting a [pull request](#) to our Evals repo with examples of the task.

# Deprecation history

All deprecations are listed below, with the most recent announcements at the top.

## 2024-10-02: Assistants API beta v1

In [April 2024](#) when we released the v2 beta version of the Assistants API, we announced that access to the v1 beta would be shut off by the end of 2024. Access to the v1 beta will be discontinued on December 18, 2024.

See the Assistants API v2 beta [migration guide](#) to learn more about how to migrate your tool usage to the latest version of the Assistants API.

SHUTDOWN DATE	MODEL / SYSTEM	RECOMMENDED REPLACEMENT
2024-12-18	OpenAI-Beta: assistants=v1	OpenAI-Beta: assistants=v2

## 2024-08-29: Fine-tuning training on babbage-002 and davinci-002 models

On August 29th, 2024, we notified developers fine-tuning `babbage-002` and `davinci-002` that new fine-tuning training runs on these models will no longer be supported starting October 28, 2024.

Fine-tuned models created from these base models are not affected by this deprecation, but you will no longer be able to create new fine-tuned versions with these models.

SHUTDOWN DATE	MODEL / SYSTEM	RECOMMENDED REPLACEMENT
2024-10-28	New fine-tuning training on babbage-002	<code>gpt-4o-mini</code>
2024-10-28	New fine-tuning training on davinci-002	<code>gpt-4o-mini</code>

## 2024-06-06: GPT-4-32K and Vision Preview models

On June 6th, 2024, we notified developers using `gpt-4-32k` and `gpt-4-vision-preview` of their upcoming deprecations in one year and six months respectively. As of June 17, 2024, only existing users of these models will be able to continue using them.

SHUTDOWN DATE	DEPRECATED MODEL	DEPRECATED MODEL PRICE	RECOMMENDED REPLACEMENT
2025-06-06	<code>gpt-4-32k</code>	\$60.00 / 1M input tokens + \$120 / 1M output tokens	<code>gpt-4o</code>

SHUTDOWN DATE	DEPRECATED MODEL	DEPRECATED MODEL PRICE	RECOMMENDED REPLACEMENT
2025-06-06	gpt-4-32k-0613	\$60.00 / 1M input tokens + \$120 / 1M output tokens	gpt-4o
2025-06-06	gpt-4-32k-0314	\$60.00 / 1M input tokens + \$120 / 1M output tokens	gpt-4o
2024-12-06	gpt-4-vision-preview	\$10.00 / 1M input tokens + \$30 / 1M output tokens	gpt-4o
2024-12-06	gpt-4-1106-vision-preview	\$10.00 / 1M input tokens + \$30 / 1M output tokens	gpt-4o

## 2023-11-06: Chat model updates

On November 6th, 2023, we [announced](#) the release of an updated GPT-3.5-Turbo model (which now comes by default with 16k context) along with deprecation of `gpt-3.5-turbo-0613` and `gpt-3.5-turbo-16k-0613`. As of June 17, 2024, only existing users of these models will be able to continue using them.

SHUTDOWN DATE	DEPRECATED MODEL	DEPRECATED MODEL PRICE	RECOMMENDED REPLACEMENT
2024-09-13	gpt-3.5-turbo-0613	\$1.50 / 1M input tokens + \$2.00 / 1M output tokens	gpt-3.5-turbo
2024-09-13	gpt-3.5-turbo-16k-0613	\$3.00 / 1M input tokens + \$4.00 / 1M output tokens	gpt-3.5-turbo

Fine-tuned models created from these base models are not affected by this deprecation, but you will no longer be able to create new fine-tuned versions with these models.

## 2023-08-22: Fine-tunes endpoint

On August 22nd, 2023, we [announced](#) the new fine-tuning API (`/v1/fine_tuning/jobs`) and that the original `/v1/fine-tunes` API along with legacy models (including those fine-tuned with the `/v1/fine-tunes` API) will be shut down on January 04, 2024. This means that models fine-tuned using the `/v1/fine-tunes` API will no longer be accessible and you would have to fine-tune new models with the updated endpoint and associated base models.

### Fine-tunes endpoint

SHUTDOWN DATE	SYSTEM	RECOMMENDED REPLACEMENT
2024-01-04	<code>/v1/fine-tunes</code>	<code>/v1/fine_tuning/jobs</code>

## 2023-07-06: GPT and embeddings

On July 06, 2023, we [announced](#) the upcoming retirements of older GPT-3 and GPT-3.5 models served via the completions endpoint. We also announced the upcoming retirement of our first-generation text embedding models. They will be shut down on January 04, 2024.

**InstructGPT models**

SHUTDOWN DATE	DEPRECATED MODEL	DEPRECATED MODEL PRICE	RECOMMENDED REPLACEMENT
2024-01-04	text-ada-001	\$0.40 / 1M tokens	gpt-3.5-turbo-instruct
2024-01-04	text-babbage-001	\$0.50 / 1M tokens	gpt-3.5-turbo-instruct
2024-01-04	text-curie-001	\$2.00 / 1M tokens	gpt-3.5-turbo-instruct
2024-01-04	text-davinci-001	\$20.00 / 1M tokens	gpt-3.5-turbo-instruct
2024-01-04	text-davinci-002	\$20.00 / 1M tokens	gpt-3.5-turbo-instruct
2024-01-04	text-davinci-003	\$20.00 / 1M tokens	gpt-3.5-turbo-instruct

Pricing for the replacement `gpt-3.5-turbo-instruct` model can be found on the [pricing page](#).

**Base GPT models**

SHUTDOWN DATE	DEPRECATED MODEL	DEPRECATED MODEL PRICE	RECOMMENDED REPLACEMENT
2024-01-04	ada	\$0.40 / 1M tokens	babbage-002
2024-01-04	babbage	\$0.50 / 1M tokens	babbage-002
2024-01-04	curie	\$2.00 / 1M tokens	davinci-002
2024-01-04	davinci	\$20.00 / 1M tokens	davinci-002
2024-01-04	code-davinci-002	---	gpt-3.5-turbo-instruct

Pricing for the replacement `babbage-002` and `davinci-002` models can be found on the [pricing page](#).

**Edit models & endpoint**

SHUTDOWN DATE	MODEL / SYSTEM	RECOMMENDED REPLACEMENT
2024-01-04	text-davinci-edit-001	gpt-4o
2024-01-04	code-davinci-edit-001	gpt-4o
2024-01-04	/v1/edits	/v1/chat/completions

**Fine-tuning GPT models**

SHUTDOWN DATE	DEPRECATED MODEL	TRAINING PRICE	USAGE PRICE	RECOMMENDED REPLACEMENT
2024-01-04	ada	\$0.40 / 1M tokens	\$1.60 / 1M tokens	babbage-002
2024-01-04	babbage	\$0.60 / 1M tokens	\$2.40 / 1M tokens	babbage-002
2024-01-04	curie	\$3.00 / 1M tokens	\$12.00 / 1M tokens	davinci-002
2024-01-04	davinci	\$30.00 / 1M tokens	\$120.00 / 1K tokens	davinci-002, gpt-3.5-turbo, gpt-4o

### First-generation text embedding models

SHUTDOWN DATE	DEPRECATED MODEL	DEPRECATED MODEL PRICE	RECOMMENDED REPLACEMENT
2024-01-04	text-similarity-ada-001	\$4.00 / 1M tokens	text-embedding-3-small
2024-01-04	text-search-ada-doc-001	\$4.00 / 1M tokens	text-embedding-3-small
2024-01-04	text-search-ada-query-001	\$4.00 / 1M tokens	text-embedding-3-small
2024-01-04	code-search-ada-code-001	\$4.00 / 1M tokens	text-embedding-3-small
2024-01-04	code-search-ada-text-001	\$4.00 / 1M tokens	text-embedding-3-small
2024-01-04	text-similarity-babbage-001	\$5.00 / 1M tokens	text-embedding-3-small
2024-01-04	text-search-babbage-doc-001	\$5.00 / 1M tokens	text-embedding-3-small
2024-01-04	text-search-babbage-query-001	\$5.00 / 1M tokens	text-embedding-3-small
2024-01-04	code-search-babbage-code-001	\$5.00 / 1M tokens	text-embedding-3-small
2024-01-04	code-search-babbage-text-001	\$5.00 / 1M tokens	text-embedding-3-small
2024-01-04	text-similarity-curie-001	\$20.00 / 1M tokens	text-embedding-3-small
2024-01-04	text-search-curie-doc-001	\$20.00 / 1M tokens	text-embedding-3-small
2024-01-04	text-search-curie-query-001	\$20.00 / 1M tokens	text-embedding-3-small
2024-01-04	text-similarity-davinci-001	\$200.00 / 1M tokens	text-embedding-3-small
2024-01-04	text-search-davinci-doc-001	\$200.00 / 1M tokens	text-embedding-3-small
2024-01-04	text-search-davinci-query-001	\$200.00 / 1M tokens	text-embedding-3-small

## 2023-06-13: Updated chat models

On June 13, 2023, we announced new chat model versions in the [Function calling and other API updates](#) blog post. The three original versions will be retired in June 2024 at the earliest. As of January 10, 2024, only existing users of these models will be able to continue using them.

SHUTDOWN DATE	LEGACY MODEL	LEGACY MODEL PRICE	RECOMMENDED REPLACEMENT
at earliest 2024-06-13	gpt-4-0314	\$30.00 / 1M input tokens + \$60.00 / 1M output tokens	gpt-4o

SHUTDOWN DATE	DEPRECATED MODEL	DEPRECATED MODEL PRICE	RECOMMENDED REPLACEMENT
2024-09-13	gpt-3.5-turbo-0301	\$15.00 / 1M input tokens + \$20.00 / 1M output tokens	gpt-3.5-turbo
2025-06-06	gpt-4-32k-0314	\$60.00 / 1M input tokens + \$120.00 / 1M output tokens	gpt-4o

## 2023-03-20: Codex models

SHUTDOWN DATE	DEPRECATED MODEL	RECOMMENDED REPLACEMENT
2023-03-23	code-davinci-002	gpt-4o
2023-03-23	code-davinci-001	gpt-4o
2023-03-23	code-cushman-002	gpt-4o
2023-03-23	code-cushman-001	gpt-4o

## 2022-06-03: Legacy endpoints

SHUTDOWN DATE	SYSTEM	RECOMMENDED REPLACEMENT
2022-12-03	/v1/engines	<a href="#">/v1/models</a>
2022-12-03	/v1/search	<a href="#">View transition guide</a>
2022-12-03	/v1/classifications	<a href="#">View transition guide</a>
2022-12-03	/v1/answers	<a href="#">View transition guide</a>