

# Pricing

## Latest models

Text tokens

Price per 1M tokens · Batch API price ☐

Model	Input	Cached input	Output
<div>gpt-4o</div> <div>↳ gpt-4o-2024-08-06</div>	\$2.50	\$1.25	\$10.00
<div>gpt-4o-audio-preview</div> <div>↳ gpt-4o-audio-preview-2024-12-17</div>	\$2.50	-	\$10.00
<div>gpt-4o-realtime-preview</div> <div>↳ gpt-4o-realtime-preview-2024-12-17</div>	\$5.00	\$2.50	\$20.00
<div>gpt-4o-mini</div> <div>↳ gpt-4o-mini-2024-07-18</div>	\$0.15	\$0.075	\$0.60
<div>gpt-4o-mini-audio-preview</div> <div>↳ gpt-4o-mini-audio-preview-2024-12-17</div>	\$0.15	-	\$0.60
<div>gpt-4o-mini-realtime-preview</div> <div>↳ gpt-4o-mini-realtime-preview-2024-12-17</div>	\$0.60	\$0.30	\$2.40
<div>o1</div> <div>↳ o1-2024-12-17</div>	\$15.00	\$7.50	\$60.00
<div>o3-mini</div> <div>↳ o3-mini-2025-01-31</div>	\$1.10	\$0.55	\$4.40
<div>o1-mini</div> <div>↳ o1-mini-2024-09-12</div>	\$1.10	\$0.55	\$4.40

▼ All snapshots

Audio tokens

Price per 1M tokens

Model	Input	Cached input	Output
<div>gpt-4o-audio-preview</div> <div>↳ gpt-4o-audio-preview-2024-12-17</div>	\$40.00	-	\$80.00
<div>gpt-4o-mini-audio-preview</div> <div>↳ gpt-4o-mini-audio-preview-2024-12-17</div>	\$10.00	-	\$20.00
<div>gpt-4o-realtime-preview</div> <div>↳ gpt-4o-realtime-preview-2024-12-17</div>	\$40.00	\$2.50	\$80.00

Model	Input	Cached input	Output
<div>gpt-4o-mini-realtime-preview</div> <div>↳ gpt-4o-mini-realtime-preview-2024-12-17</div>	\$10.00	\$0.30	\$20.00

▼ All snapshots

## Fine tuning

Price per 1M tokens · Batch API price ☐

Model	Training	Input	Cached input	Output
gpt-4o-2024-08-06	\$25.00	\$3.75	\$1.875	\$15.00
gpt-4o-mini-2024-07-18	\$3.00	\$0.30	\$0.15	\$1.20
gpt-3.5-turbo	\$8.00	\$3.00	-	\$6.00
davinci-002	\$6.00	\$12.00	-	\$12.00
babbage-002	\$0.40	\$1.60	-	\$1.60

## Assistants API

The tokens used for the Assistants API are billed at the chosen model's per-token rates.

GB refers to binary gigabytes of storage (also known as gibibyte), where 1GB is 2^30 bytes.

Tool	Cost
Code Interpreter	\$0.03 session
File Search	\$0.10 GB/day (1GB free)

## Transcription and speech generation

Model	Use case	Cost
Whisper	Transcription	\$0.006 minute
TTS	Speech generation	\$15.00 1M characters
TTS HD	Speech generation	\$30.00 1M characters

## Image generation

Price per image

Model	Quality	1024x1024	1024x1792
DALL·E 3	Standard	\$0.04	\$0.08

Model	Quality	1024x1024	1024x1792
	HD	\$0.08	\$0.12
Model	256x256	512x512	1024x1024
DALL·E 2	\$0.016	\$0.018	\$0.02

## Embeddings

Price per 1M tokens · Batch API price ☐

Model	Cost
text-embedding-3-small	\$0.02
text-embedding-3-large	\$0.13
text-embedding-ada-002	\$0.10

## Moderation

Model	Cost
omni-moderation-latest ↳ omni-moderation-2024-09-26	Free
text-moderation-latest ↳ text-moderation-007	Free

▼ All snapshots

## Other models

Price per 1M tokens · Batch API price ☐

Model	Input	Output
chatgpt-4o-latest	\$5.00	\$15.00
gpt-4-turbo ↳ gpt-4-turbo-2024-04-09	\$10.00	\$30.00
gpt-4 ↳ gpt-4-0613	\$30.00	\$60.00
gpt-4-32k	\$60.00	\$120.00
gpt-3.5-turbo ↳ gpt-3.5-turbo-0125	\$0.50	\$1.50
gpt-3.5-turbo-instruct	\$1.50	\$2.00
gpt-3.5-turbo-16k-0613	\$3.00	\$4.00
davinci-002	\$2.00	\$2.00
babbage-002	\$0.40	\$0.40

✓ All snapshots